

# 虛擬化技術 *Virtualization Technique*

## System Virtualization

Memory Virtualization



國立清華大學

# *Agenda*

- Basic concept of memory management
- Brief introduction to ARM v7 AMSA
- Memory Virtualization
  - Shadow page table
  - Hardware assistance
  - Comparison

A decorative blue curved shape on the left side of the slide, transitioning from a solid blue to a lighter blue gradient.

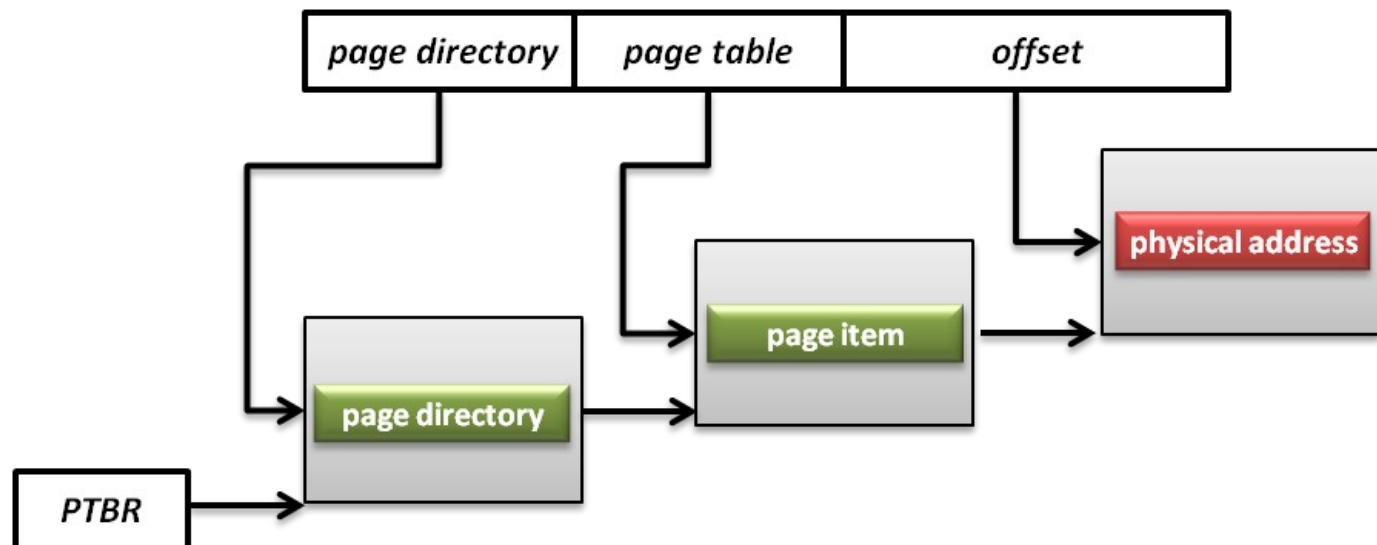
*Basic concept of memory management*

# *Memory Management Unit*

- A hardware component responsible for handling accesses to memory requested by the CPU
  - Address translation: virtual address to physical address (VA to PA)
  - Memory protection
  - Cache control
  - Bus arbitration
- Page tables are maintained by operating system, and MMU only references them.
- TLB updates are performed automatically by page-table walking hardware

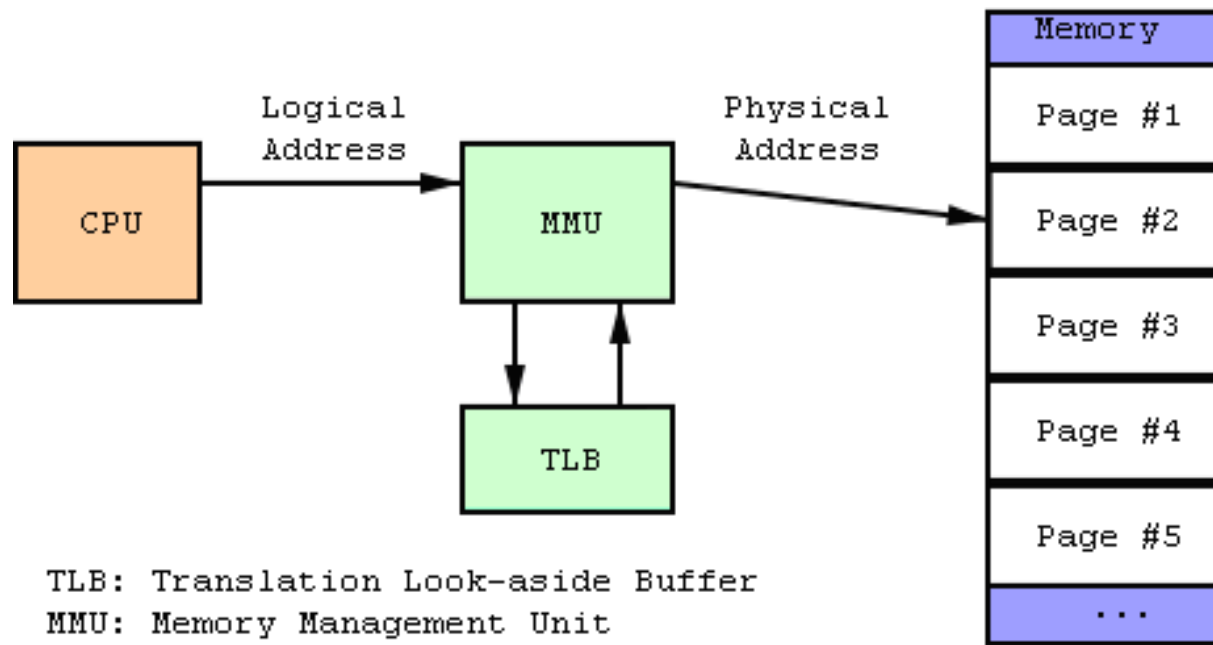
# Page Tables

- A page table is the data structure used by a virtual memory system to store the mapping between virtual addresses and physical addresses
- Translation table base register(TTBR)
  - Also called page table base register
  - A register that stores the address of the base page table for MMU



- Translation look-aside buffer

- A CPU cache that memory management hardware uses to improve virtual address translation speed
- The TLB is typically implemented as content-addressable memory (CAM)
- The CAM search key is the virtual address and the search result is a physical address



TLB: Translation Look-aside Buffer

MMU: Memory Management Unit

CPU: Central Processing Unit

A decorative blue curved graphic element on the left side of the slide, consisting of several concentric, semi-transparent arcs that create a sense of depth and movement.

## *Brief introduction to ARMv7 VMSA*

# *ARMv7 VMSA*






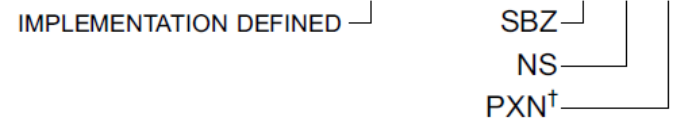
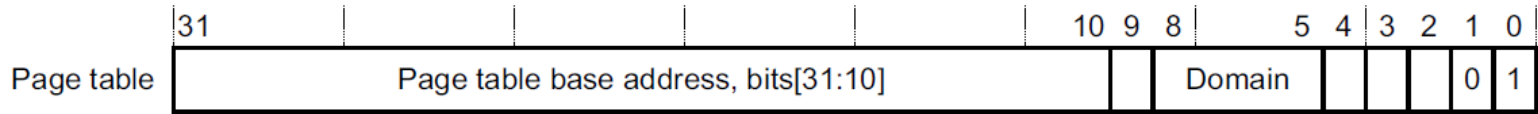

- VMSA: Virtual memory system architecture
- Support four mapping types
  - 16MB Super section
  - 1MB Section
  - 64KB Large page
  - 4KB Small page
- Basically a 32-bit two-layer translation table
- Other features (virtualization extension and 64-bit large physical address extension) would be mentioned in the future class



# *L1 Page Table*

- First-level table
  - Translation properties for a Section and Super-section
  - Translation properties and pointers to a second-level table for a Large page or a Small page.
- L1 descriptor[1:0]
  - 0b00, invalid or fault entry
    - Translation fault
  - 0b01, page table
    - Address of 2<sup>nd</sup> level translation table
  - 0b10, section of super-section
  - 0b11, special usage with PXN attribute

## 1<sup>st</sup> Page Table Descriptor



# *L2 Page Table*

- Second-level tables
  - The base address and translation properties for a Small page or a Large page.
  - That is, page tables
  - 1KB
- L2 descriptor[1:0]
  - 0b00, invalid or fault entry
    - Translation fault
  - 0b01, large page
  - 0b1x, small page

## 2<sup>nd</sup> Page Table Descriptor

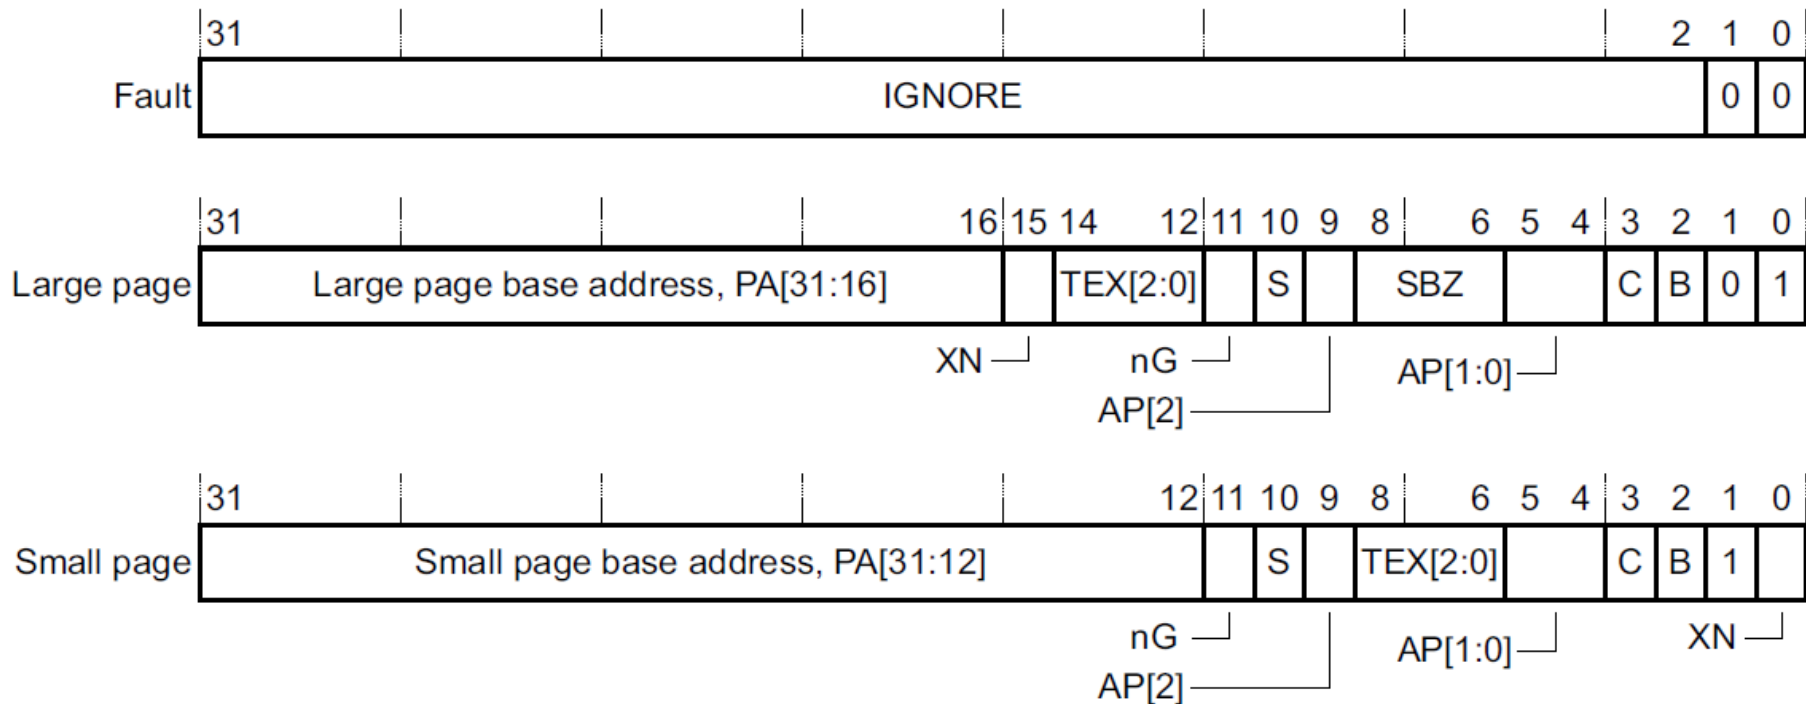


Figure B3-5 Short-descriptor second-level descriptor formats



## Concepts

Shadow page table

Hardware assistance

Comparison


# *Memory Virtualization*

# *Memory Management on a VM*

- Traditionally, OS fully controls all physical memory space and provides a continuous addressing space to each process
- Guest OS is just one of user space processes of host OS
- If guest OS is allowed to access the physical memory arbitrarily, then what happens?
- In system virtualization, VMM should make all virtual machines share the physical memory space


# Memory Virtualization

- Memory virtualization architecture



# Memory Virtualization

- The performance drop of memory access is usually unbearable. VMM needs further optimization.
- VMM maintains shadow page tables :
  - Direct virtual-to-physical address mapping
  - Use hardware TLB for address translation





# *Goals of Memory Virtualization*

- Address Translation
  - Control table-walking hardware that accesses translation tables in main memory.
- Memory Protection
  - Define access permission which uses the Access Control Hardware.
- Access Attribute
  - Define attribute and type of memory region to direct how memory operation to be handled.
- How to implement?
  - Software solution: shadow page table
  - Hardware solution
    - NPT on SVM from AMD
    - EPT on VMX from Intel
    - ARM v7 VMSA (Virtual Memory System Architecture) with virtualization extension



Concepts

Shadow page table

Hardware assistance


Comparison

## *Memory Virtualization*

# *Shadow Page Table*

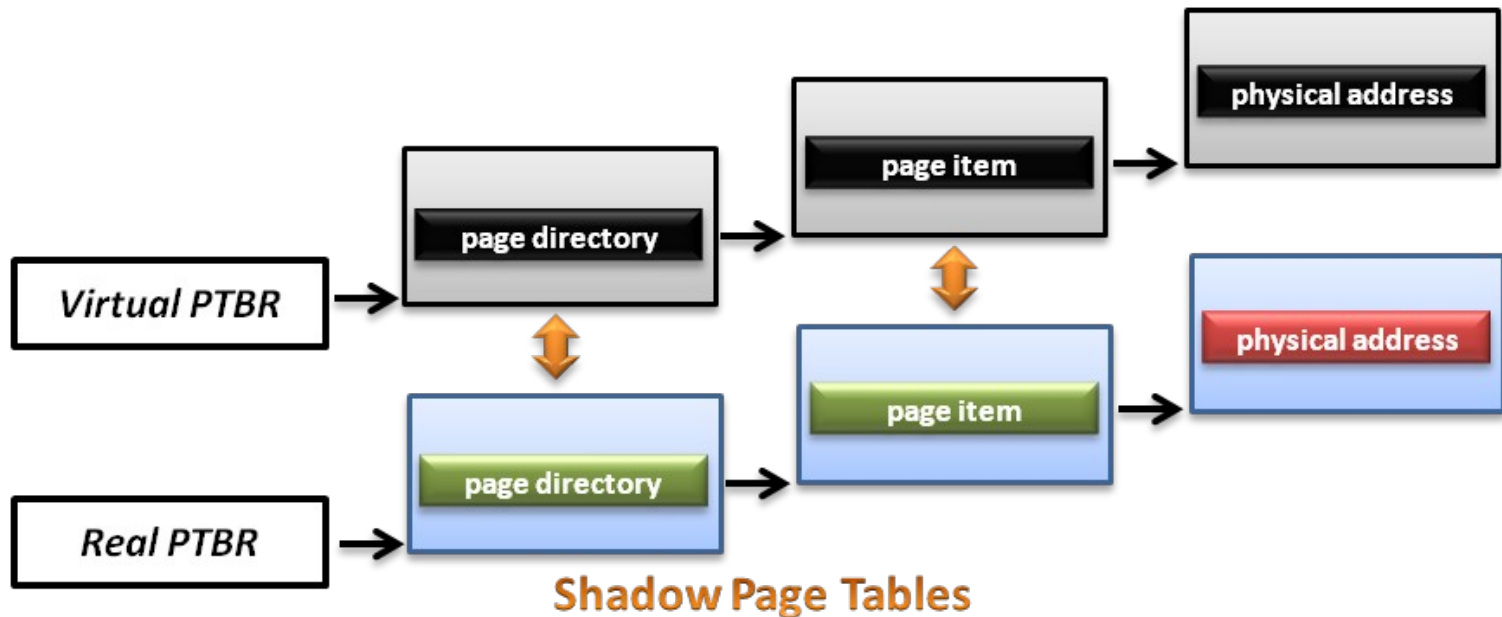
- Map guest virtual address to host physical address
  - Shadow page table
    - Guest OS will maintain its own virtual memory page table in the guest physical memory frames.
    - For each guest physical memory frame, VMM should map it to host physical memory frame.
    - Shadow page table maintains the mapping from guest virtual address to host physical address.
  - Page table protection
    - VMM will apply write protection to all the physical frames of guest page tables, which lead the guest page table write exception and trap to VMM.

# Shadow Page Table: Overview




# Shadow Page Table

- How does this technique work ?
  - VMM should make MMU virtualized
    - VMM manages the real PTBR and a virtual PTBR for each VM
    - When a guest OS is activated, the real PTBR points to the corresponding shadow page table of the guest OS
    - When the guest OS attempts to modify the PTBR, it will be intercepted by VMM for further emulation



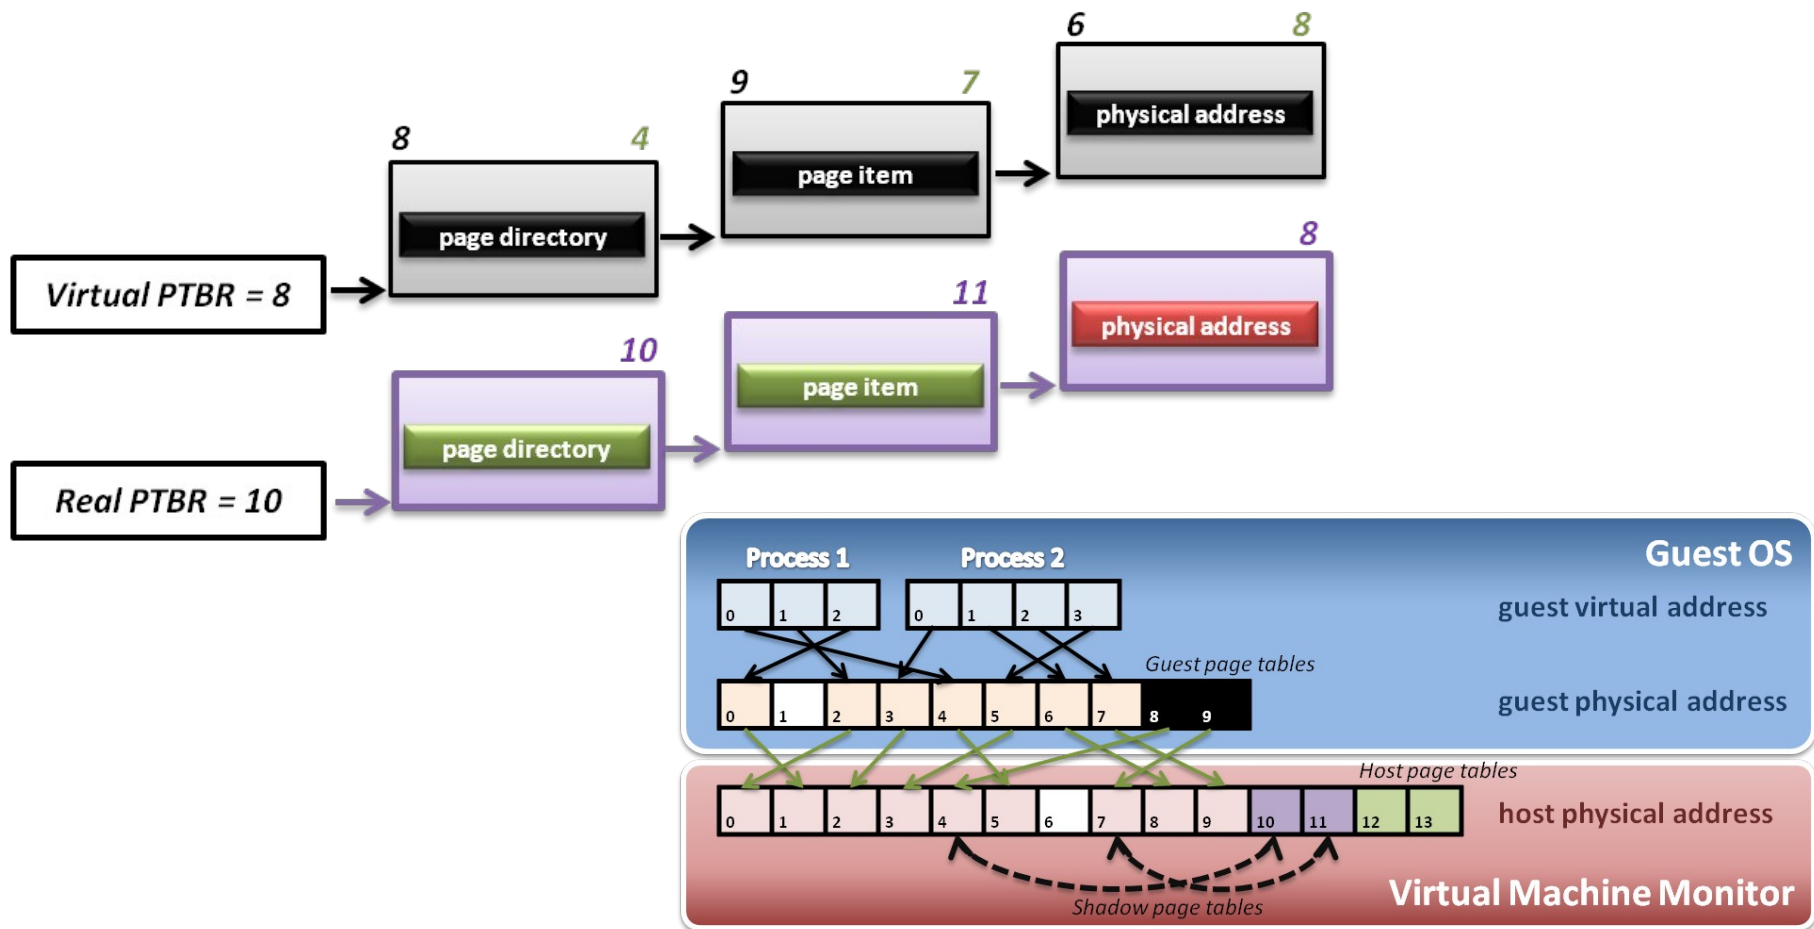
# Shadow Page Table

- Construct shadow page table
  - Guest OS will maintain its own page table for each process.
  - VMM maps each guest physical page to host physical page.
  - Create shadow page tables for each guest page table.
  - VMM should protect host frame which contains guest page table. (that means to write protect the guest page tables in host memory)



# Shadow Page Table

- Implement with PTBR :
  - For example, process 2 in guest OS wants to access its memory whose page number is 1.





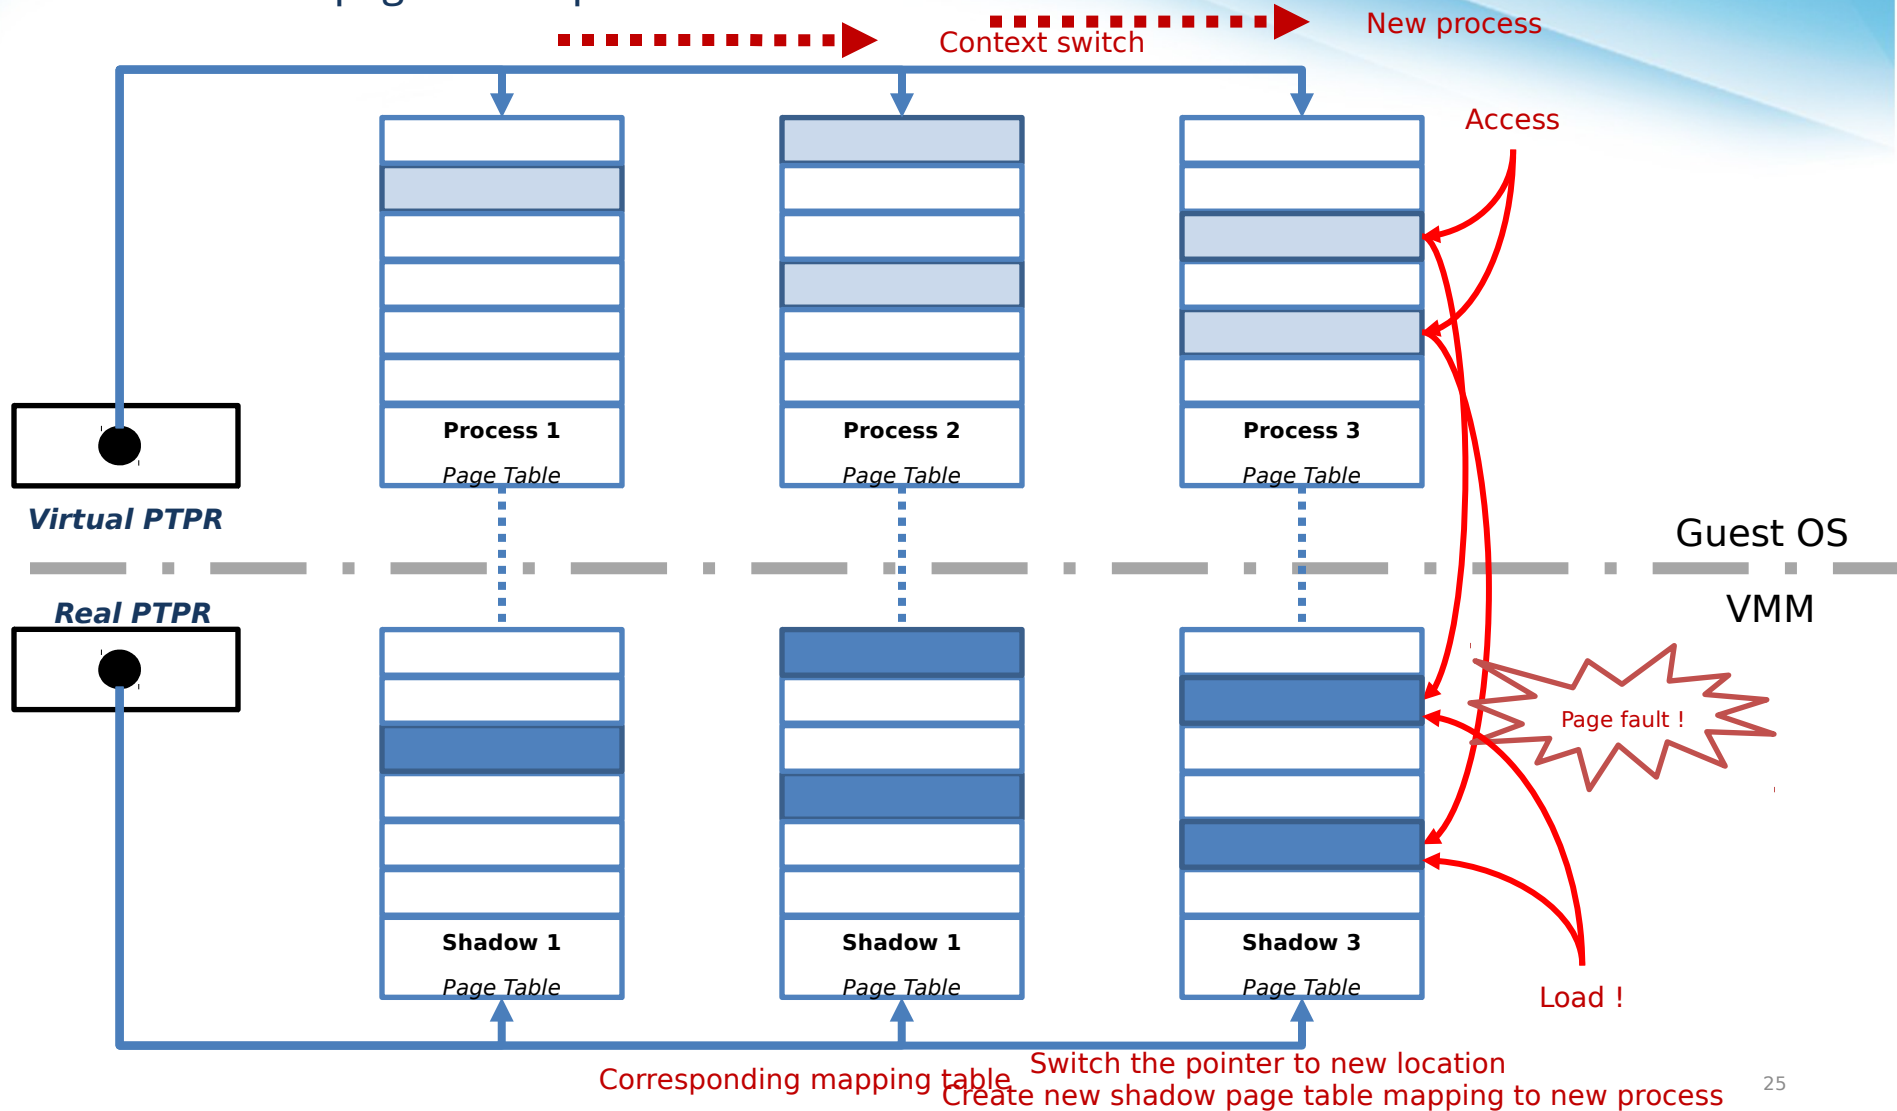
# *SPT Maintenance*

- If guest OS modify one of its page tables, then the corresponding entry of SPT must be updated.
  - We call it “shadow” because SPT is just like the shadow of page tables of guest OS.
- How to identify this kind of modification?
  - Guest OS could read/write a physical frame with the help of SPT.
  - Mark those physical frames used as guest page tables read-only, so that when a guest OS tries to modify its guest page table, an exception would be triggered.
  - Then VMM checks the modification and updates the corresponding entry on SPT



# Shadow Page Table

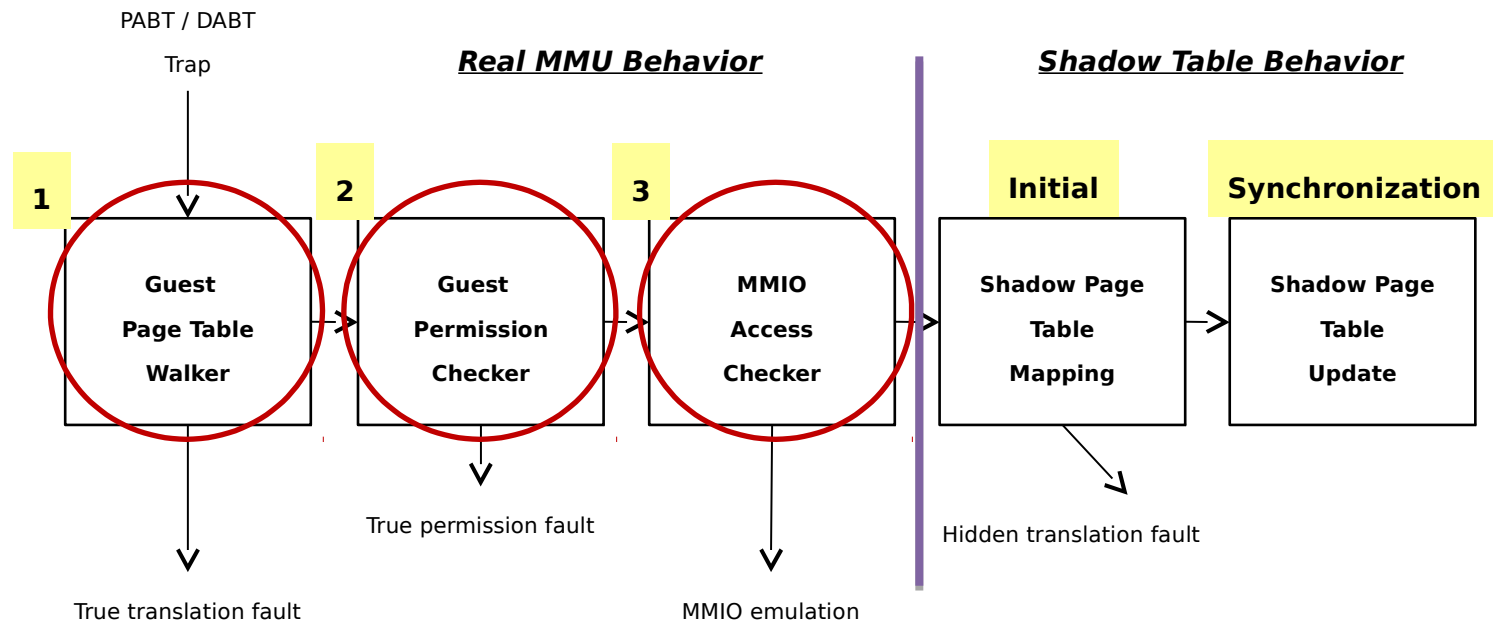
- Shadow page table operations :



# *Big Overhead*

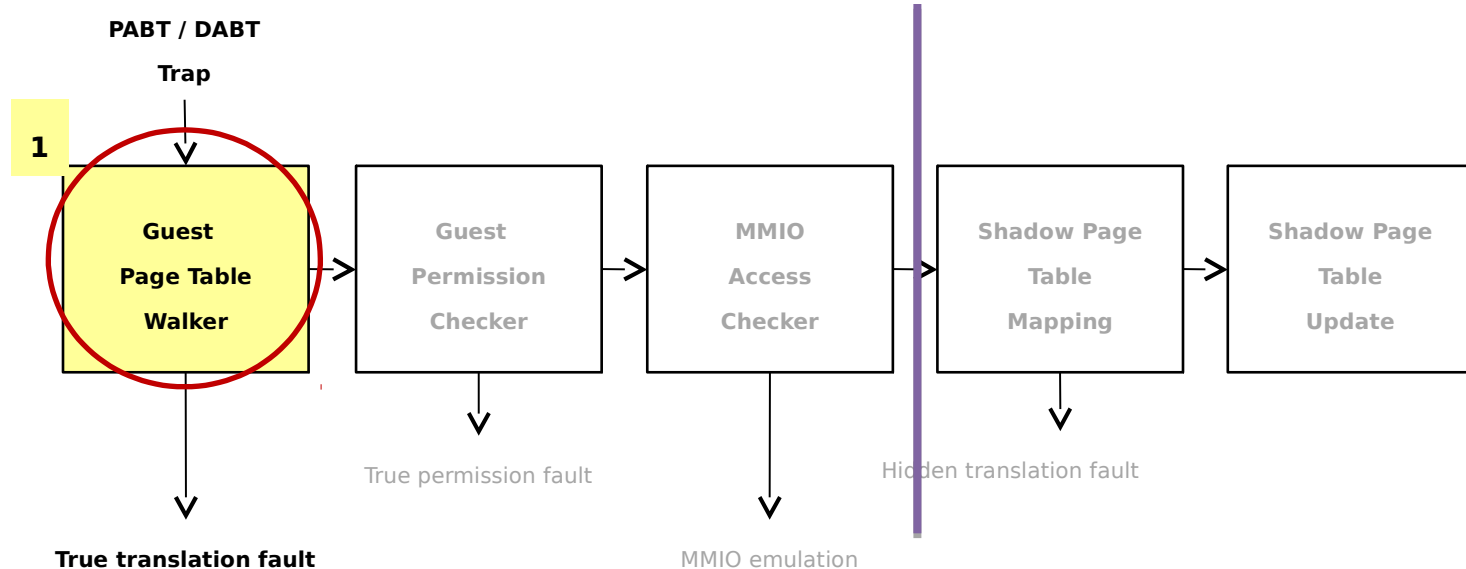
- A page fault caused by guest OS would launch the walking process that costs a lot of overhead.
  - Several steps to get a new entry on shadow page table
    - Walk page tables on guest OS
    - Check the permission on guest
    - Offset shift: GPA to HVA
    - Walk page tables on VMM
    - Check the permission on VMM
    - New entry established!
    - Invalidate the TLB entry
  - Each new process on guest OS would consume two pages. One is the page table on guest OS, and the other is the corresponding shadow page table.

# Page Walking Process on ARM



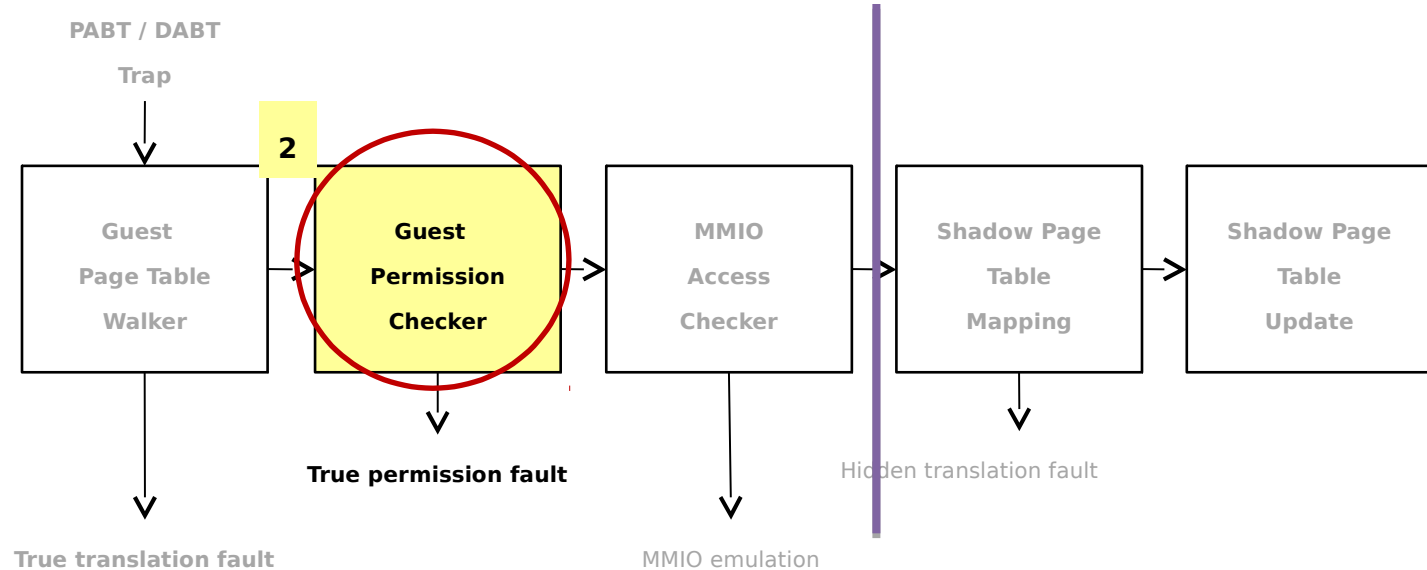
# Step 1

- While a page fault is occurred, Guest Page Table Walker will walk through guest page table to check if the fault is from guest



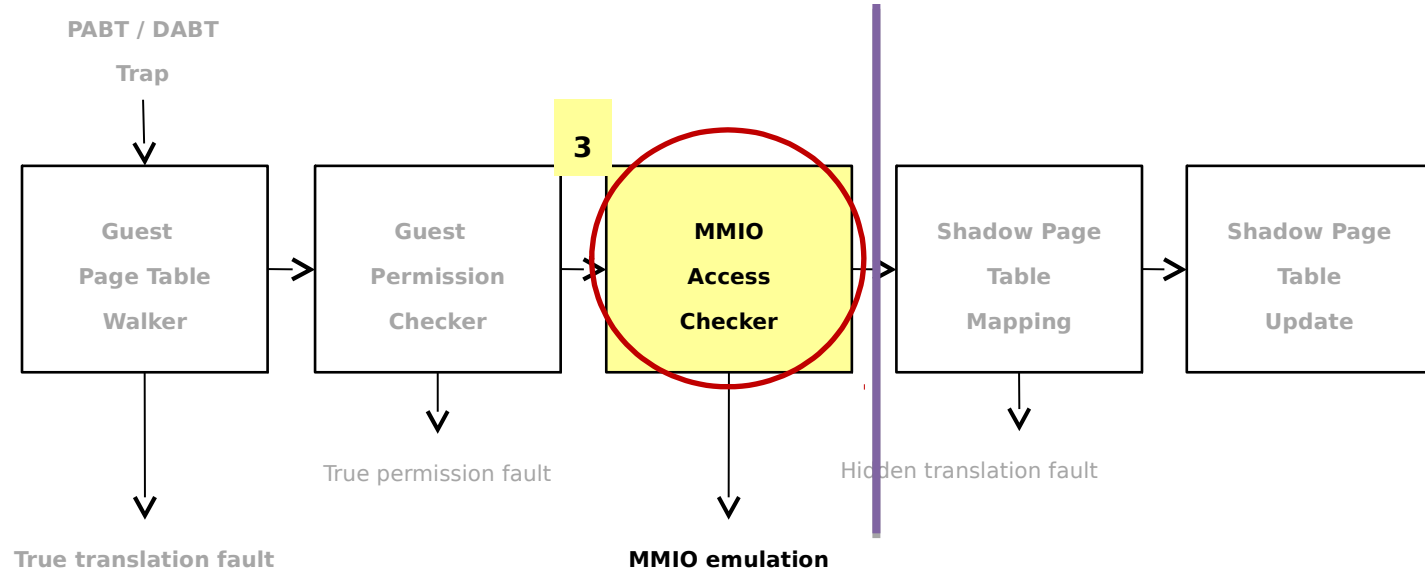
## Step 2

- Step 2 will check if guest access permission is not allowed



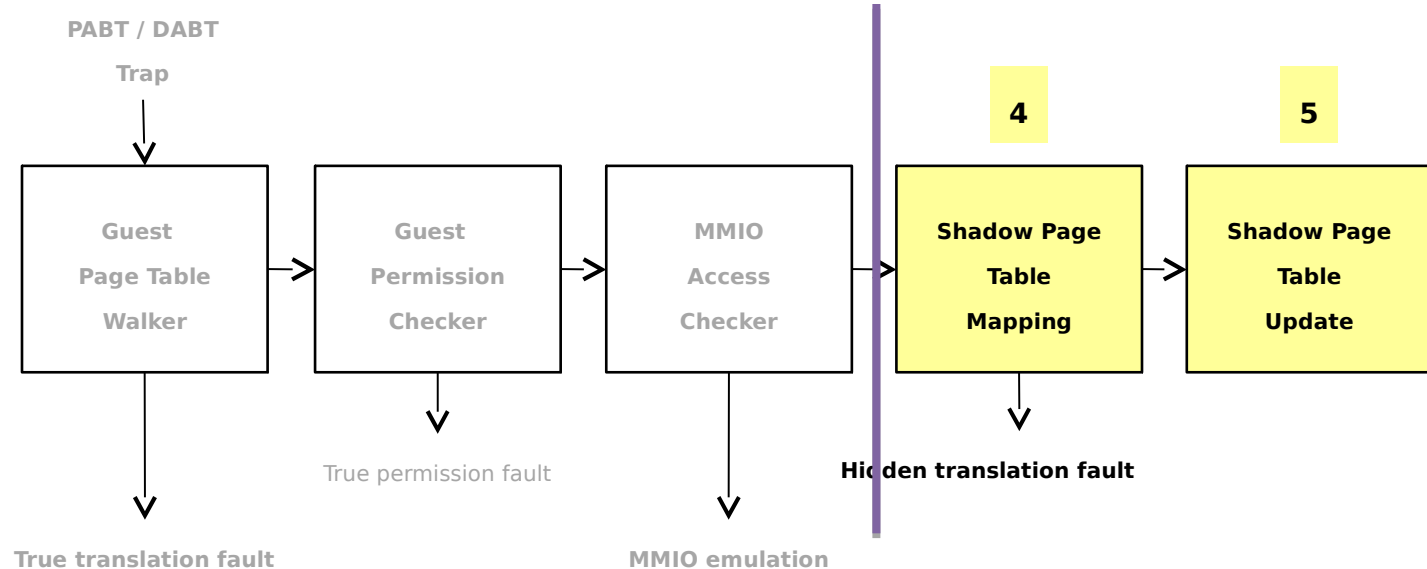
## Step 3

- Step 3 will check if the guest physical memory address used is located in the range of MMIO address



## Steps 4 & 5

- Step 4 and step 5 are used to build up shadow page tables and maintain their consistency between guest and shadow ones



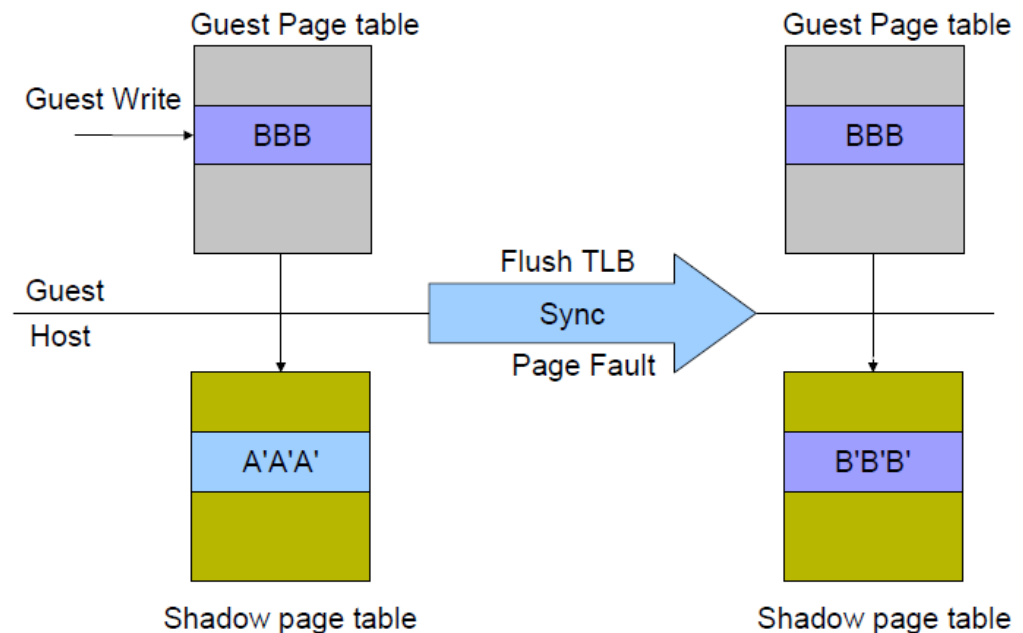
# *Optimization For SPT*

- Para-virtualization
  - Reduce VM exits
  - Guest OS would send a hyper call to VMM when guest OS sets the page table entries.
  - This method will eliminate from using write-protection for synchronization
- No trap for non-present PTEs
  - Reduce VM exits
  - If the page-fault is caused by PTE not present, it is not intercepted by the host
  - Only works on VMX





# Optimization For SPT

- Un-synchronize shadow page table pages
  - Reduce VM exits
  - Allow the guest page table to be writable if and only if the page is the last level page-structure (level 1)
  - Based on TLB rules
    - We need to flush TLB to ensure the translation use the modified page structures, then we can Intercept the TLB flush operations and sync shadow pages
    - Sometimes, TLB need not be flushed, then it can be synced through page fault



# Optimization for SPT

- KSM: kernel shared memory



*The shadow tables of kernel space are shared by all guest processes*

## *Other Issues*

- Page fault and page protection issue
  - When a physical page fault occurs, VMM needs to decide whether this exception should be injected to guest OS or not
    - If the page entry in a guest page table is still valid, VMM prepares for the corresponding page and does not inject any exception to guest OS.
    - If the page entry in a guest page table is invalid, then VMM directly injects the virtual page fault to guest OS.



Concepts

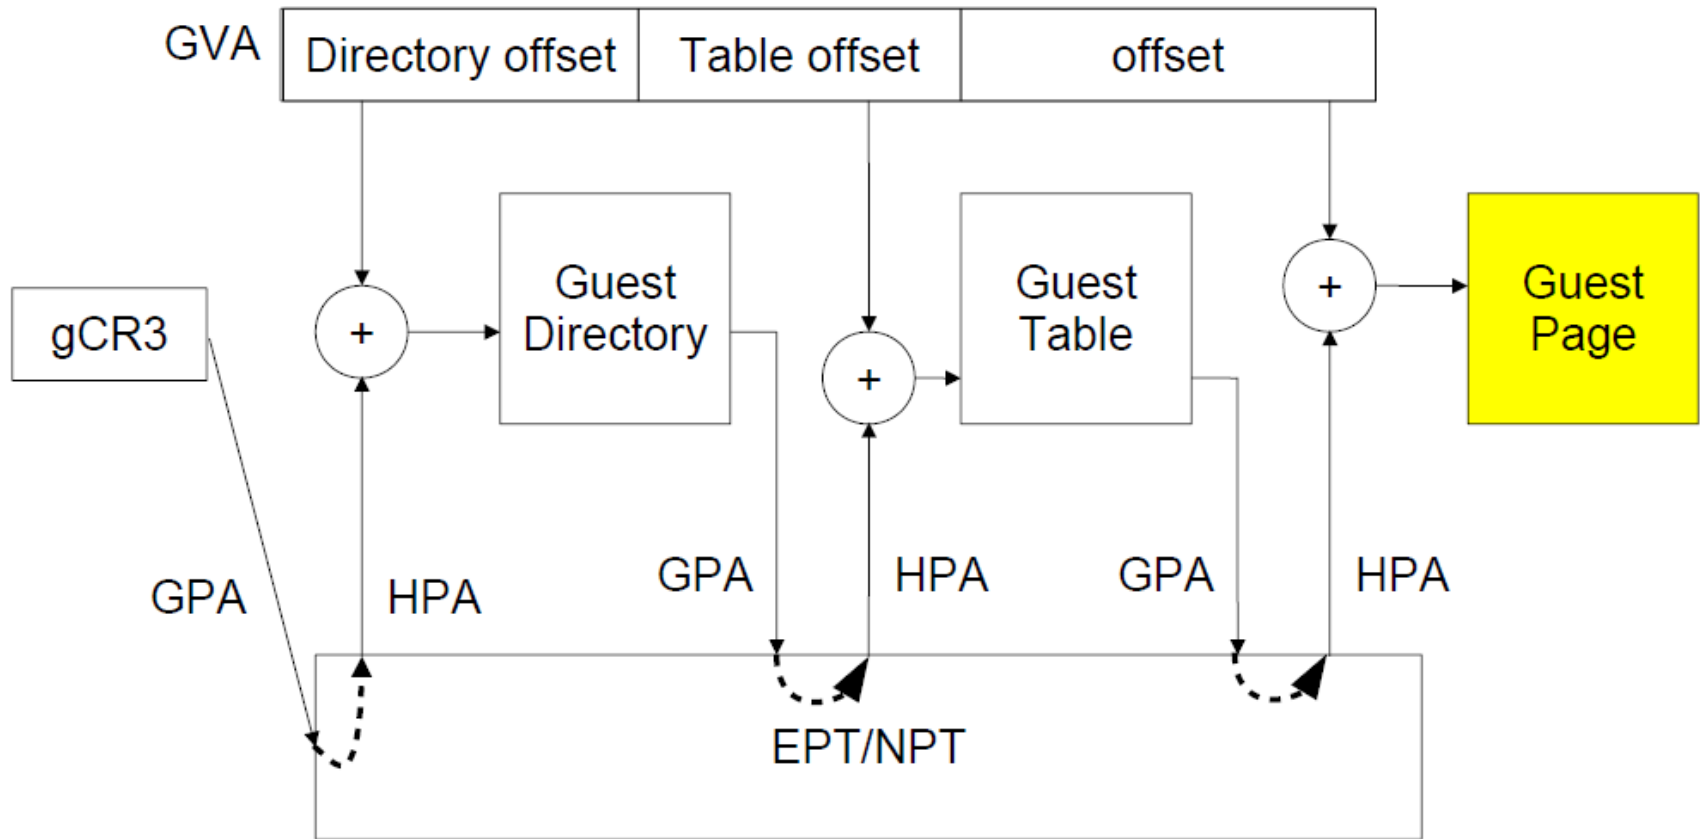
Shadow page table

Hardware assistance

Comparison

## ***Memory Virtualization***

# Hardware Assistance: Overview



# *Hardware Solution*

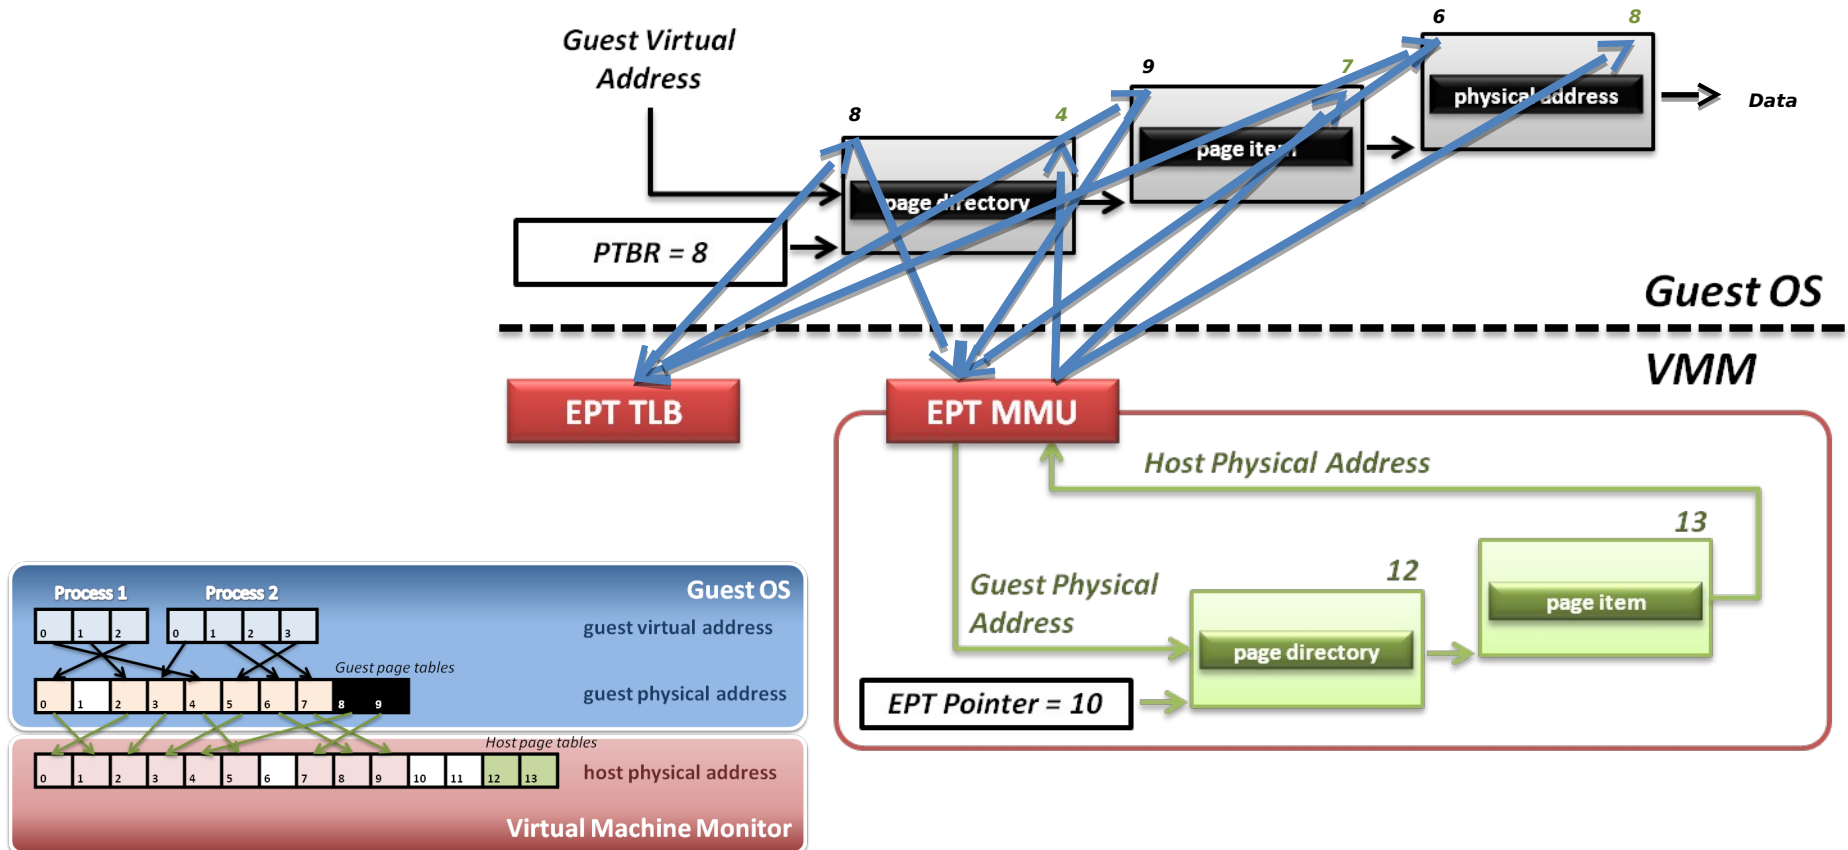
- Difficulties of shadow page table technique :
  - Shadow page table implementation is extremely complex.
  - Page fault mechanism and synchronization issues are critical.
  - Host memory space overhead is considerable.
- But why we need this technique to virtualize MMU ?
  - MMU do not first implemented for virtualization.
  - MMU is knowing nothing about two level page address translation.
- Now, let us consider hardware solution.

# *Extended Page Table*

- Concept of Extended Page Table (EPT) :
  - Instead of walking along with only one page table hierarchy, EPT technique implement one more page table hierarchy.
    - One page table is maintained by guest OS, which is used to generate guest physical address.
    - The other page table is maintained by VMM, which is used to map guest physical address to host physical address.
  - For each memory access operation, EPT MMU directly gets the guest physical address from guest page table, and then gets the host physical address by the VMM mapping table automatically.

# Extended Page Table

- Memory operation :







Concepts

Shadow page table

Hardware assistance

Comparison

## ***Memory Virtualization***

## *Question*

- Computer architecture with virtualization extension is a trend.
- Hardware-assisted techniques replace many software methods of virtualization.
- However, is hardware-assisted implementation a definite winner?

# *Comparisons*

## Hardware-assisted

- Walk any requested address
  - Appropriate to programs that have a large amount of page table miss when executing
  - Less chance to exit VM (less context switch)
- Two-layer EPT
  - Means each access needs to walk two tables
- Easier to develop
  - Many particular registers
  - Hardware helps guest OS to notify the VMM

## Software solution


- Only walk when SPT entry miss
  - Appropriate to programs that would access only some addresses frequently
  - Every access might be intercepted by VMM (many traps)
- One reference
  - Fast and convenient when page hit
- Hard to develop
  - Two-layer structure
  - Complicated reverse map
  - Permission emulation

# *Combination?*

- Is hardware-assisted implementation a definite winner?
  - No
  - How about combining these two methods?
- Selective memory virtualization
  - The VMM can dynamically choose the memory management mechanism depending on the executing status.
  - Challenge
    - Hard to figure out the standard that we judge a program performs too many page table misses
    - An accurate algorithm to sample behaviors of a program is necessary
  - Group from Peking University completed a selective solution on VEE 2011
    - Gain just a little bit performance improvement
    - For now, it's not worth to do so.

# Selective Memory Virtualization

- Experiment data
  - The bar shows the normalized execution time
  - The lower, the better.
  - HAP: hardware-assisted page table
  - SP: shadow page table
  - DSP: dynamic selective page table
  - Take hardware extension as 100%
- We can see that...
  - Generally HAP is faster than SPT, but not always.
  - The performance of DSP is best, but only tiny disparities. (about 2%)



# *Memory Virtualization Summary*

- Software implementation

- Memory architecture

- MMU (memory management unit)
    - TLB (translation look-aside buffer)

- Shadow page table

- MMU virtualization by virtual PTBR
    - Shadow page table construction
    - Page fault and page table protection

- Hardware assistance

- Extended page table

- Hardware walk guest and host page table simultaneously

# *Reference*

- Selective hardware/software memory virtualization
  - <http://www.cs.mtu.edu/~zlwang/papers/vee11.pdf>
- ARM<sup>®</sup> Architecture Reference Manual: ARMv7-A and ARMv7-R edition